

Optimizing the software stack of a cosmic proportions cluster of multi-core machines

Sebastian Pop

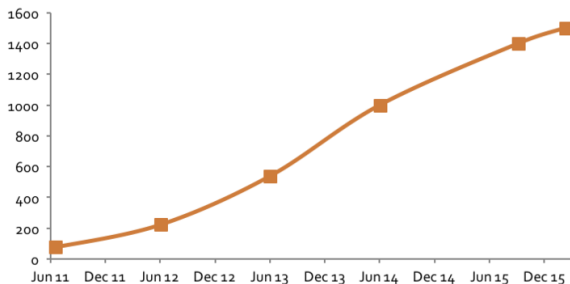
SARC: Samsung Austin R&D Center

February 5, 2017

Android: a cosmic size cluster

- ▶ top500: 10M cores / 15.3MW¹ / US\$273 million²
- ▶ Android devices: $\sim 6B$ cores³ / $\sim 300MW$ ⁴ / \sim US\$0

Google Android active base (m)



[Source: Google, a16z]

¹<https://www.top500.org/lists/2016/11>

²https://en.wikipedia.org/wiki/Sunway_TaihuLight

³4 cores / device

⁴battery 13.2Wh = 4.4V * 3000mAh, charging every 48 hours

Android Open Source Project (AOSP) Software Stack

- ▶ AOSP: common base for Android devices (+ customizations)
- ▶ C/C++ for the platform libraries, Java for user interface

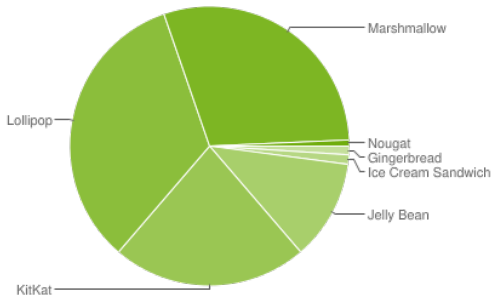
ansic	22 MLoC	39%
cpp	13 MLoC	23%
java	10 MLoC	17%
- ▶ ~ 80% execution cycles in C/C++, ~ 20% in Java

⁵Data collected during a 7-day period ending on January 9, 2017.

Android Open Source Project (AOSP) Software Stack

- ▶ AOSP: common base for Android devices (+ customizations)
- ▶ C/C++ for the platform libraries, Java for user interface

ansic	22 MLoC	39%
cpp	13 MLoC	23%
java	10 MLoC	17%
- ▶ ~ 80% execution cycles in C/C++, ~ 20% in Java
- ▶ release/updates/deprecation (5 ~ 6 years)



Why Optimizing the Performance of Android?

Why bothering?

- ▶ the code of Android is cold (flat profile), full of branches
- ▶ there are few loops (image processing, compression, etc.)

⁶\$0.12/kWh, battery 13.2Wh = 4.4V * 3000mAh, charging every 48 hours

Why Optimizing the Performance of Android?

Why bothering?

- ▶ the code of Android is cold (flat profile), full of branches
- ▶ there are few loops (image processing, compression, etc.)

Motivation:

- ▶ same code executed billions of time
- ▶ outer loop is outside the device
- ▶ profile how often code is in use
- ▶ variation over time following popularity of apps
- ▶ continuously monitor usage patterns
- ▶ tune code optimizations over time

\$0.30 / device / year \rightarrow \$300M / billion devices / year⁶



⁶\$0.12/kWh, battery 13.2Wh = 4.4V * 3000mAh, charging every 48 hours

Agenda

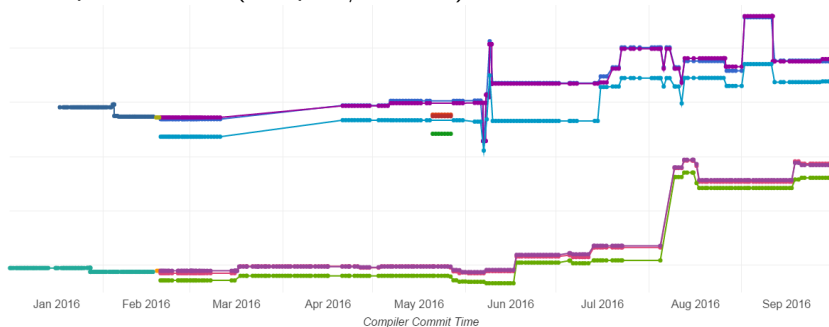
- ▶ Performance analysis: hot spots
- ▶ Improve performance of AOSP libraries
- ▶ Enable continuous profiling and optimizations (AutoFDO)
- ▶ Enable more secure execution environments (CFI)

Performance Analysis

- ▶ benchmarks: track performance over time (compiler/libraries)
- ▶ linux perf: profile of cycles (per function, hot-spots)
- ▶ valgrind: number of executed instructions (branches, R/W)
- ▶ static profiles: how many uses for a function

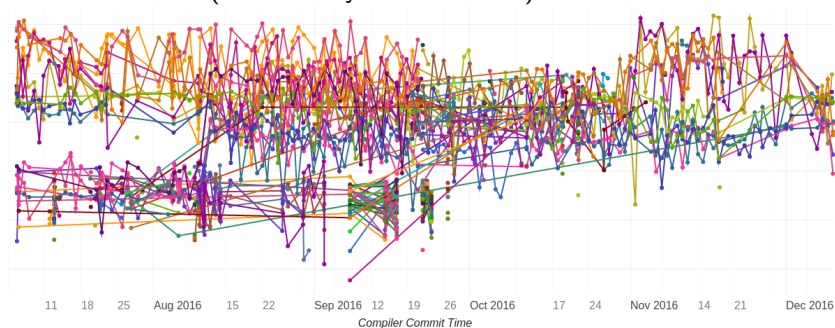
Benchmarks

track performance (compiler/libraries) over time



Benchmarks

on a real device (and a noisy benchmark...)



Linux Perf

Valgrind

Static Profile

- ▶ -fllto: static call-graph, estimated frequencies per call

Improve performance of AOSP libraries

SARC contributions

- ▶ update Android NDK libc++, make it easy to keep updated
- ▶ 20x speedup of `std::string.find()` in libc++ and libstdc++
need to port perf to memmem and strstr of bionic and glibc
- ▶ improve perf of `shared_ptr` in libc++
- ▶ improve perf of string to int value parsing in libc++

Benchmarking Standard Libraries


SARC contribution: std-benchmark⁷

- ▶ std-benchmark provides micro-benchmarks for functions in libc and C++ standard library
- ▶ detect room for improvement
 - ▶ compile with different compilers
 - ▶ link with different standard libraries
 - ▶ run on different machines: CPUs, architectures

⁷<https://github.com/hiraditya/std-benchmark>

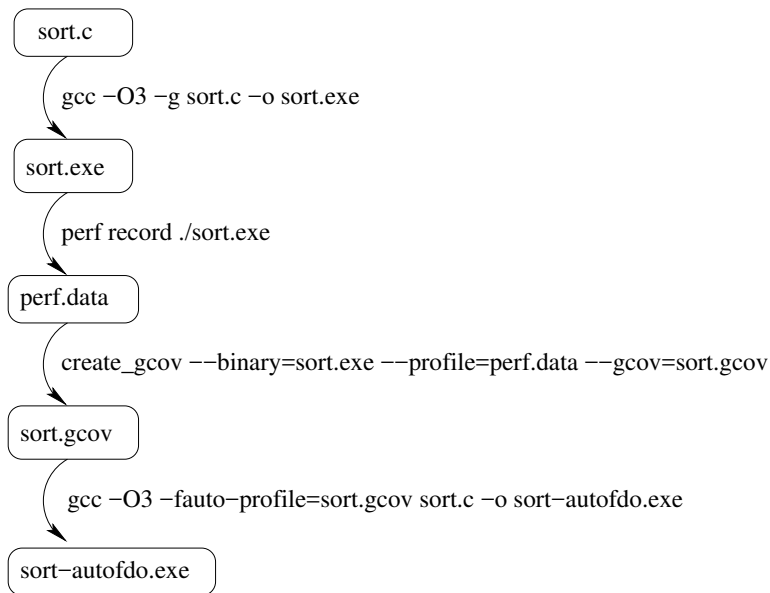
AutoFDO: Feedback Directed Optimization

- ▶ linux-perf extracts profiles of running systems
- ▶ little to no overhead ⁸
- ▶ coverage (basic block frequencies) from dynamic profiles
- ▶ continuous profiling and tuning of optimizations



⁸Google Wide Profiling: A Continuous Profiling Infrastructure for Data Centers, IEEE Micro (2010)

AutoFDO: Example



AutoFDO: Code Optimizations

- ▶ better inlining ⁹, devirtualization, function instantiation
- ▶ hot/cold code placement
- ▶ register allocation, jump-threading, etc.

⁹Lightweight Feedback-Directed Cross-Module Optimization, CGO 2010

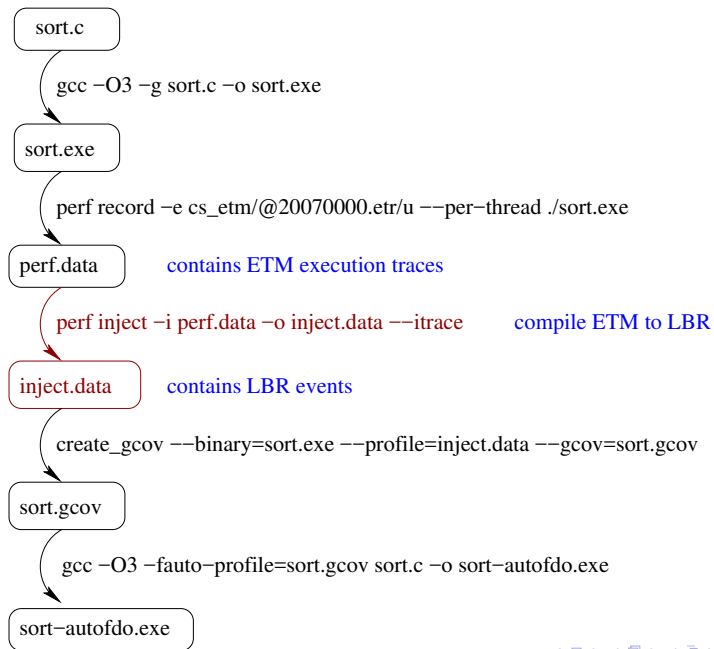
AutoFDO: More Precise Coverage

- ▶ Intel-LBR (Last Branch Record): last 16 taken branches
- ▶ provides more precise basic block execution frequency
- ▶ how do we do this on ARM?

ARM-ETM: Embedded Trace Macrocell

- ▶ ARM-ETM: records execution traces (for debug)
 - ▶ dedicated circular buffer 1 to 3MB ($\sim 10^5$ branches/MB)
 - ▶ no overhead
 - ▶ support in Linux kernel by Mathieu Poirier (Linaro)
 - ▶ next android kernel linux-4.9 will support ARM-ETM
-
- ▶ **SARC contribution:** how to use ARM-ETM for AutoFDO
 - ▶ perf-inject translates execution traces to LBR events
 - ▶ patch similar to perf-inject for Intel Process Trace

AutoFDO: with ARM-ETM



From Dynamic Profiles to Power Usage

- ▶ traditionally, per app battery usage (ammeter on wire) ¹⁰
- ▶ more accurate picture with linux-perf profiles:
 - ▶ profiles from the field: real world use-cases
 - ▶ merge together different profiles
 - ▶ compute code execution frequency
 - ▶ power consumption estimation per line of code

¹⁰An Analysis of Power Consumption in a Smartphone, USENIX'10

Towards more secure devices

- ▶ Control Flow Integrity (CFI): 2% overhead ¹¹
- ▶ to enable on Android: need to further reduce its cost

¹¹Enforcing Forward-Edge Control-Flow Integrity in GCC&LLVM, USENIX'14